

ICP Assignment-4

Nirmala Yarlagadda
700733102

GitHub link: <https://github.com/niryarjessy22/ICP-4.git>

Videos link:

https://drive.google.com/file/d/1sAZAFFEnqPxIOXqLx-QSkQIXBEpQAHSg/view?usp=share_link

1. Data Manipulation

a. Read the provided CSV file 'data.csv'.

```
In [1]: #Importing all the required Libraries
import pandas as pd
import numpy as np
import nltk
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #Reading the CSV file 'data.csv'.
data=pd.read_csv('data.csv')
```

```
In [3]: #Viewing the shape of the data
data.shape
```

```
Out[3]: (169, 4)
```

```
In [4]: #Getting the information of the data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 169 entries, 0 to 168
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   Duration    169 non-null    int64   
 1   Pulse       169 non-null    int64   
 2   Maxpulse    169 non-null    int64   
 3   Calories    164 non-null    float64  
dtypes: float64(1), int64(3)
memory usage: 5.4 KB
```

```
In [5]: #Show the basic statistical description about the data.
data.describe()
```

```
Out[5]:
```

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

```
In [6]: #Checking for the null values in the dataset
data.isnull().any().sum()
```

```
Out[6]: 1
```

- d. Check if the data has null values.
- e. Select at least two columns and aggregate the data using: min, max, count, mean.
- f. Filter the dataframe to select the rows with calories values between 500 and 1000.
- i. Replace the null values with the mean

```
In [7]: #Replace the null values with the mean
df = data.fillna(data.mean())
```

```
In [8]: #this step shows the number of null values is zero after replaing with mean value
df.isnull().any().sum()
```

```
Out[8]: 0
```

```
In [9]: #agregating the data using using: mean
df[['Duration', 'Pulse']].mean()
```

```
Out[9]: Duration    63.846154
Pulse          107.461538
dtype: float64
```

```
In [10]: #agregating the data using using: min
df[['Duration', 'Pulse']].min()
```

```
Out[10]: Duration    15
Pulse             80
dtype: int64
```

```
In [11]: #agregating the data using using: max,
df[['Duration', 'Pulse']].max()
```

```
Out[11]: Duration    300
Pulse             159
dtype: int64
```

```
In [12]: #agregating the data using using:count
df[['Duration', 'Pulse']].count()
```

```
Out[12]: Duration    169
Pulse             169
dtype: int64
```

g. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.

```
In [13]: #Filtering the dataframe to select the rows with calories values between 500 and 1000.  
print(df[(df['Calories'] < 1000) & (df['Calories'] > 500)])
```

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

```
In [14]: #Filter the dataframe to select the rows with calories values > 500 and pulse < 100.  
print(df[(df['Pulse'] < 100) & (df['Calories'] > 500)])
```

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

```
In [15]: #Creating a new "df_modified" dataframe that contains all the columns from df except for maxpulse  
df_modified= df.drop(['Maxpulse'], axis=1)
```

```
In [16]: df_modified.shape
```

```
Out[16]: (169, 3)
```

```
In [17]: #the dataset after dropping the maxpulse  
df_modified.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 169 entries, 0 to 168  
Data columns (total 3 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Duration    169 non-null    int64  
1   Pulse       169 non-null    int64  
2   Calories    169 non-null    float64  
dtypes: float64(1), int64(2)  
memory usage: 4.1 KB
```

- h. Create a new “df_modified” dataframe that contains all the columns from df except for “Maxpulse”.
- i. Delete the “Maxpulse” column from the main df dataframe
- j. Convert the datatype of Calories column to int datatype.

```
In [18]: #Delete the "Maxpulse" column from the main df dataframe  
del df["Maxpulse"]
```

```
In [19]: #the dataset after delating the maxpulse  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 169 entries, 0 to 168  
Data columns (total 3 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Duration    169 non-null    int64  
1   Pulse       169 non-null    int64  
2   Calories    169 non-null    float64  
dtypes: float64(1), int64(2)  
memory usage: 4.1 KB
```

```
In [20]: #Converting the datatype of Calories column to int datatype.  
df = df.astype({"Calories":'int'})
```

```
In [21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 169 entries, 0 to 168  
Data columns (total 3 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Duration    169 non-null    int64  
1   Pulse       169 non-null    int64  
2   Calories    169 non-null    int32  
dtypes: int32(1), int64(2)  
memory usage: 3.4 KB
```

k. Using pandas create a scatter plot for the two columns (Duration and Calories).
Example

```
In [24]: #Scatter plot for Duration against Calories
plt.scatter(df['Duration'], data['Calories'])

#Displaying the title for the plot
plt.title("plot for Curation vs calories")

# Setting the X and Y Labels
plt.xlabel('Duration')
plt.ylabel('Calories')

plt.show()
```



2. Linear Regression

- Import the given "Salary_Data.csv"
- Split the data in train_test partitions, such that 1/3 of the data is reserved as test subset.
- Train and predict the model.
- Calculate the mean_squared error
- Visualize both train and test data using scatter plot.

```
In [36]: #Importing all the required libraries
import pandas as pd
import numpy as np
import nltk
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
from sklearn.metrics import mean_squared_error

from sklearn.linear_model import LinearRegression
```

```
In [37]: #Reading the the given "Salary_Data.csv"
data=pd.read_csv('Salary_Data.csv')
```

```
In [38]: #basic info about the dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```

```
In [39]: X = data.iloc[:, :-1].values
Y = data.iloc[:, 1].values
```

```
In [40]: # Splitting the dataset into the Training set and Test set
X_Trainingset, X_Testingset, Y_Trainingset, Y_Testingset = train_test_split(X, Y, test_size=0.3, random_state=0)
```

```
In [41]: # Fitting Simple Linear Regression to the training set
regressor = LinearRegression()
regressor.fit(X_Trainingset, Y_Trainingset)

# Predicting the values
Y_Pred = regressor.predict(X_Testingset)

# calculate mean_square error
mse = mean_squared_error(Y_Testingset,Y_Pred)
print(f"\nMean Square Error = {mse}")
```

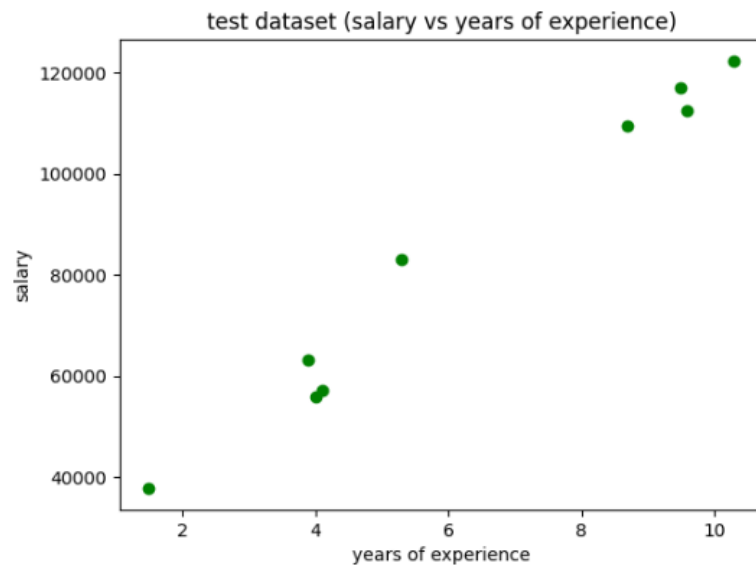
```
Mean Square Error = 23370078.800832972
```

visualize both train and test data using scatter plot


```
In [42]: # Visualising the Training set results
plt.scatter(X_Trainingset, Y_Trainingset, color='green')
plt.title('training dataset (salary vs years of experience)')
plt.xlabel('years of experience')
plt.ylabel('salary')
plt.show()
```



```
In [43]: # Visualising the Test set results
plt.scatter(X_Testingset, Y_Testingset, color='green')
plt.title('test dataset (salary vs years of experience)')
plt.xlabel('years of experience')
plt.ylabel('salary')
plt.show()
```



```
In [44]: # plotting the actual and predicted values
compare = [i for i in range(1, len(Y_Testingset)+1, 1)]
plt.plot(compare, Y_Testingset, color='blue', linestyle='--')
plt.plot(compare, Y_Pred, color='green', linestyle='--')
plt.xlabel('index')
plt.ylabel('salary')
plt.title('actual salary value(blue line) vs Predicted Salary values(green line)')
plt.show()
```




```
In [45]: # Plotting the Final Output i.e., the test data and predicted data
plt.scatter(X_Testingset, Y_Testingset, color='green')
plt.plot(X_Testingset, Y_Pred, color='black', linewidth=3)
plt.title('salary vs years of experience')
plt.xlabel('years of experience')
plt.ylabel('salary')
plt.show()
```



GitHub link: <https://github.com/niryarjessy22/ICP-4.git>

Videos link:

https://drive.google.com/file/d/1sAZAFFEnqPxIOXqLx-QSkQIXBEpQAHSG/view?usp=share_link